

DAT565/DIT407 Assignment 6

Saif Sayed
gussayedfa@student.gu.se

Gona Ibrahim Abdulrahman
gusibrigo@student.gu.se

2024-05-10

This is a report for assignment 6 for the course *Introduction to Data Science & AI* from Chalmers and Gothenburg University.

Our source code can be found in Appendix A of this document.

Problem 1: The dataset

To load and verify the MNIST dataset, the following steps were undertaken:

Setting up PyTorch: The code included a comment indicating the installation of PyTorch via conda. This step ensures that the necessary PyTorch libraries are available.

Data Loading and Transformation: A transform was defined using the *transforms.Compose* function from the torchvision library. The transform applied was *transforms.ToTensor()* which converts the images to PyTorch tensors.

Loading the Training and Test Datasets: The MNIST dataset was loaded using the *datasets.MNIST* function from torchvision. The root parameter specified the directory where the dataset should be stored. The train parameter was set to True for the training dataset and False for the test dataset. The download parameter was set to True to automatically download the dataset if it was not already present. The transform parameter was set to the defined transform to apply the transformation to the loaded images.

Plotting Images: A function named *plot_images* was defined to plot sample images from the dataset. The function utilized the *matplotlib.pyplot* library to create a figure and axes for plotting the images. The function iterated over a range of indices and retrieved the corresponding image from the dataset. The image was then plotted using the *imshow* function, with the title indicating whether it belonged to the training or test dataset. The images were displayed using the *plt.show()* function.

Image Dimensions and Value Scale Verification: A loop was implemented to iterate over the training dataset. Within each iteration, an image and its corresponding label were retrieved from the dataset. Assertions were used to verify

that the image dimensions were (1, 28, 28), indicating a grayscale image with a size of 28x28 pixels. Another assertion ensured that the pixel values of the image were normalized to a range between 0 and 1. If any of the assertions failed, an appropriate error message was displayed.

Results: The code successfully loaded the MNIST dataset and verified the image dimensions and value scale for the training dataset. The plot of sample images from both the training and test datasets provided a visual representation of the dataset (see **Figure 1** and **2**).




Figure 1: Visual Representation of the MNIST dataset (training set).




Figure 2: Visual Representation of the MNIST dataset (test set).

Problem 2: Single hidden layer

To solve the problem, the following steps were undertaken:

Data Loading and Batch Training: Data loaders were created using the `torch.utils.data.DataLoader` function. The training and test datasets were passed into the data loaders, along with the batch size and shuffle parameters. The training data loader shuffles the data during training to introduce randomness.

Device Selection: The code checks if a GPU is available using `torch.cuda.is_available()`. If a GPU is available, the device is set to 'cuda'; otherwise, it is set to 'cpu'.

Network Architecture: A custom PyTorch module called Net was defined to represent the feedforward neural network. The module was initialized with the specified number of input units, hidden units, and output classes. The network architecture includes a linear layer, followed by the ReLU activation function, and another linear layer.

Model Training: The model was instantiated using the defined network architecture and moved to the selected device. The loss function, `nn.CrossEntropyLoss()`, was defined to calculate the cross-entropy loss. The optimizer, `optim.SGD`, was initialized with the model parameters and the specified learning rate. A loop was implemented to iterate over the specified number of epochs. Within each epoch, the model was set to training mode, and the training dataset was iterated over in batches. The forward pass was performed, and the loss was calculated. The gradients were set to zero, and the backward pass was executed to compute the gradients. The optimizer was used to update the model parameters.

Model Evaluation: After each epoch, the model was set to evaluation mode. The test dataset was used to calculate the accuracy of the model predictions. The number of correct predictions and the total number of predictions were accumulated. The validation accuracy was calculated by dividing the number of correct predictions by the total number of predictions and multiplying by 100. The validation accuracy was printed for monitoring the model's performance.

Results: The training process was performed for the specified number of epochs, and the validation accuracy was reported after each epoch (see Table 1).

Table 1: Validation Accuracy for single hidden layer.

Epoch	Validation Accuracy (%)
1	92.56
2	94.64
3	95.35
4	96.05
5	96.62
6	96.78
7	96.94
8	96.83
9	97.46
10	97.36

Problem 3: Two hidden layers

To solve the problem, the following steps were undertaken:

Parameters: Relevant parameters such as the number of units in the hidden layers, the number of output classes, learning rate, weight decay for L2 regularization, number of epochs, and the device (GPU or CPU) were defined.

Data Loading and Transformation: The MNIST dataset, which consists of handwritten digit images, was loaded using torchvision. Transformations were applied to convert the images to tensors.

Neural Network Model: A custom PyTorch module called Net was defined to represent the fully-connected feedforward network. The model architecture included two hidden layers with ReLU activation functions. The forward method was implemented to specify the forward pass of the network.

Loss Function and Optimizer: The cross-entropy loss function was selected for multi-class classification tasks. The SGD optimizer was used to update the model parameters. L2 regularization was enabled by setting an appropriate weight decay parameter.

Training the Model: A loop was set up to iterate over the specified number of epochs. Within each epoch, the model was set to training mode, and the training dataset was iterated over in batches. The forward pass was performed, and the loss was calculated. The gradients were set to zero, and the backward pass was executed to compute the gradients. The optimizer was used to update the model parameters.

Validation and Reporting: After each epoch, the model was set to evaluation mode. The validation dataset was used to calculate the accuracy of the model predictions. The validation accuracy was printed for monitoring the model's performance.

Results: The training process was performed for the specified number of epochs, and the validation accuracy was reported after each epoch (**see Table 2**). The goal of achieving a validation accuracy of at least 98% was successfully accomplished.

Table 2: Validation Accuracy for two hidden layers.

Epoch	Validation Accuracy (%)	Epoch	Validation Accuracy (%)
1	91.88	21	98.24
2	95.77	22	98.31
3	96.70	23	98.41
4	97.17	24	98.30
5	96.95	25	98.23
6	96.82	26	98.33
7	97.93	27	98.30
8	97.80	28	98.36
9	97.83	29	98.35
10	98.08	30	98.33
11	98.16	31	98.37
12	98.06	32	98.34
13	98.15	33	98.38
14	97.63	34	98.33
15	98.28	35	98.31
16	98.23	36	98.34
17	98.23	37	98.31
18	98.26	38	98.40
19	98.25	39	98.32
20	98.31	40	98.32

Problem 4: Convolutional neural network

To solve the problem, the following steps were taken:

Creating data loaders for batch training: The MNIST dataset is loaded using `torchvision.datasets.MNIST` and transformed using the transform object. Train, test, and validation datasets are created. Data loaders are created using `torch.utils.data.DataLoader`, which provide an iterable over the dataset for batch training.

Defining the network structure: The `ConvNet` class is defined as a subclass of `nn.Module`. The network architecture consists of two convolutional layers with ReLU activations, followed by max pooling. The output is flattened and passed through a fully connected layer with 10 output units (one for each digit class).

Initializing the network: An instance of the ConvNet class is created, representing the neural network model.

Defining the loss function and optimizer: The loss function is defined as the cross-entropy loss(`nn.CrossEntropyLoss()`). The optimizer is defined as Stochastic Gradient Descent (`optim.SGD`) with a learning rate of 0.01, momentum of 0.9 and weight loss of 0.0001.

Training the network: The network is trained for 40 epochs. For each epoch, the running loss is initialized to 0. The training data is iterated over in batches using the `train_loader`. For each batch, the gradients are set to zero using `optimizer.zero_grad()` to clear any residual gradients. The inputs are passed through the network to obtain the outputs. The loss between the outputs and the labels is calculated using the defined loss function. The gradients are computed using `loss.backward()`. The optimizer updates the model parameters using `optimizer.step()`. The loss is added to the running loss. After each epoch, the validation accuracy is calculated using the `valid_loader`. The statistics, including the epoch number, average loss, and validation accuracy, are printed (see **Table 3**).

Table 3: Validation Accuracy for convolutional neural network.

Epoch	Validation Accuracy (%)	Epoch	Validation Accuracy (%)
1	97.45	21	99.73
2	98.04	22	99.74
3	98.58	23	99.68
4	98.81	24	99.86
5	98.86	25	99.86
6	98.78	26	99.41
7	99.12	27	99.84
8	99.17	28	99.86
9	99.26	29	99.90
10	99.38	30	99.91
11	99.31	31	99.77
12	99.46	32	99.90
13	99.56	33	99.96
14	99.52	34	99.89
15	99.65	35	99.97
16	99.57	36	99.94
17	99.66	37	99.93
18	99.62	38	99.94
19	99.76	39	99.91
20	99.66	40	99.96

Appendix

A Python code

```

1  # Import necessary libraries
2  import torch
3  import torchvision
4  from torchvision import transforms
5  import matplotlib.pyplot as plt
6  from torch import nn
7  from torch import optim
8  import numpy as np
9
10 # Problem 1
11
12 # Set up PyTorch
13 # !conda install -c pytorch pytorch torchvision
    cpuonly
14
15 # Define a transform to normalize the data

```

```

16 transform = transforms.Compose([transforms.ToTensor()
17                                 ])
18 # Load the training and test datasets
19 train_data = torchvision.datasets.MNIST(root='data',
20                                         train=True, download=True, transform=transform)
21 test_data = torchvision.datasets.MNIST(root='data',
22                                         train=False, download=True, transform=transform)
23
24 # Function to plot images
25 def plot_images(dataset, title):
26     fig = plt.figure(figsize=(10, 5))
27     for i in range(10):
28         ax = fig.add_subplot(2, 5, i+1, xticks=[],
29                             yticks=[])
30         image, _ = dataset[i]
31         ax.imshow(torch.squeeze(image), cmap='gray')
32         ax.set_title(f'{title}_{i+1}')
33     plt.show()
34
35 # Plot some images from both datasets
36 plot_images(train_data, 'Train')
37 plot_images(test_data, 'Test')
38
39 # Verify image dimensions and value scale
40 for i in range(len(train_data)):
41     image, _ = train_data[i]
42     assert image.shape == (1, 28, 28), "Incorrect_
43         dimensions"
44     assert image.min() >= 0 and image.max() <= 1, "
45         Values_not_normalized"
46 print("All_images_have_correct_dimensions_and_
47     normalized_values.")
48
49 # Problem 2
50
51 # Create data loaders for batch training
52 train_loader = torch.utils.data.DataLoader(train_data,
53                                             batch_size=64, shuffle=True)
54 test_loader = torch.utils.data.DataLoader(test_data,
55                                             batch_size=64, shuffle=False)
56
57 # Set the device to GPU if available, otherwise use
58     CPU
59 device = torch.device('cuda' if torch.cuda.
60     is_available() else 'cpu')
61
62 # Define the number of hidden units in the hidden
63     layer
64 hidden_units = 100

```



```

54 # Define the number of output classes
55 num_classes = 10
56
57 # Define the learning rate
58 learning_rate = 0.001
59
60 # Define the number of epochs
61 num_epochs = 10
62
63 class Net(nn.Module):
64     def __init__(self, input_size, hidden_units,
65                 num_classes):
66         super(Net, self).__init__()
67         self.fc1 = nn.Linear(input_size, hidden_units)
68         self.relu = nn.ReLU()
69         self.fc2 = nn.Linear(hidden_units, num_classes)
70
71     def forward(self, x):
72         x = x.view(x.size(0), -1) # Flatten the input
73         images
74         x = self.fc1(x)
75         x = self.relu(x)
76         x = self.fc2(x)
77         return x
78
79 # Create an instance of the model
80 model = Net(input_size=28*28, hidden_units=
81             hidden_units, num_classes=num_classes).to(device)
82
83 criterion = nn.CrossEntropyLoss()
84 optimizer = optim.SGD(model.parameters(), lr=
85                       learning_rate)
86
87 for epoch in range(num_epochs):
88     # Set the model to training mode
89     model.train()
90
91     # Iterate over the training dataset
92     for images, labels in train_loader:
93         images = images.to(device)
94         labels = labels.to(device)
95
96         # Forward pass
97         outputs = model(images)
98         loss = criterion(outputs, labels)
99
100        # Backward and optimize
101        optimizer.zero_grad()
102        loss.backward()

```

```

99         optimizer.step()
100
101     # Set the model to evaluation mode
102     model.eval()
103
104     # Calculate the validation accuracy after each
105     epoch
106     correct = 0
107     total = 0
108     with torch.no_grad():
109         for images, labels in test_loader:
110             images = images.to(device)
111             labels = labels.to(device)
112             outputs = model(images)
113             _, predicted = torch.max(outputs.data, 1)
114             total += labels.size(0)
115             correct += (predicted == labels).sum().
116                 item()
117
118     accuracy = 100 * correct / total
119     print('Epoch [{}/{}], Validation Accuracy: {:.2f}%
120         '.format(epoch+1, num_epochs, accuracy))
121
122 # Problem 3
123
124 # Define the number of units in the hidden layers
125 hidden_units_1 = 500
126 hidden_units_2 = 300
127
128 # Define the number of output classes
129 num_classes = 10
130
131 # Define the learning rate
132 learning_rate = 0.1
133
134 # Define the weight decay for L2 regularization
135 weight_decay = 0.0001
136
137 # Define the number of epochs
138 num_epochs = 40
139
140 # Set the device to GPU if available, otherwise use
141 CPU
142 device = torch.device('cuda' if torch.cuda.
143     is_available() else 'cpu')
144
145 # Define the transformation to convert the images to
146 tensors
147 transform = transforms.ToTensor()

```

```

143 # Load the training dataset
144 train_dataset = torchvision.datasets.MNIST(root='./data', train=True, transform=transform, download=True)
145
146 # Load the test dataset
147 test_dataset = torchvision.datasets.MNIST(root='./data', train=False, transform=transform, download=True)
148
149 # Create data loaders for batch training
150 train_loader = torch.utils.data.DataLoader(
151     train_dataset, batch_size=64, shuffle=True)
152 test_loader = torch.utils.data.DataLoader(test_dataset,
153     batch_size=64, shuffle=False)
154
155 class Net(nn.Module):
156     def __init__(self, input_size, hidden_units_1,
157         hidden_units_2, num_classes):
158         super(Net, self).__init__()
159         self.fc1 = nn.Linear(input_size,
160             hidden_units_1)
161         self.relu1 = nn.ReLU()
162         self.fc2 = nn.Linear(hidden_units_1,
163             hidden_units_2)
164         self.relu2 = nn.ReLU()
165         self.fc3 = nn.Linear(hidden_units_2,
166             num_classes)
167
168     def forward(self, x):
169         x = x.view(x.size(0), -1) # Flatten the input
170         images
171         x = self.fc1(x)
172         x = self.relu1(x)
173         x = self.fc2(x)
174         x = self.relu2(x)
175         x = self.fc3(x)
176         return x
177
178 # Create an instance of the model
179 model = Net(input_size=28*28, hidden_units_1=
180     hidden_units_1, hidden_units_2=hidden_units_2,
181     num_classes=num_classes).to(device)
182
183 criterion = nn.CrossEntropyLoss()
184 optimizer = optim.SGD(model.parameters(), lr=
185     learning_rate, weight_decay=weight_decay)
186
187 for epoch in range(num_epochs):
188     # Set the model to training mode
189     model.train()

```

```

180
181     # Iterate over the training dataset
182     for images, labels in train_loader:
183         images = images.to(device)
184         labels = labels.to(device)
185
186         # Forward pass
187         outputs = model(images)
188         loss = criterion(outputs, labels)
189
190         # Backward and optimize
191         optimizer.zero_grad()
192         loss.backward()
193         optimizer.step()
194
195     # Set the model to evaluation mode
196     model.eval()
197
198     # Calculate the validation accuracy after each
199     epoch
200     correct = 0
201     total = 0
202     with torch.no_grad():
203         for images, labels in test_loader:
204             images = images.to(device)
205             labels = labels.to(device)
206             outputs = model(images)
207             _, predicted = torch.max(outputs.data, 1)
208             total += labels.size(0)
209             correct += (predicted == labels).sum().
210                 item()
211
212     accuracy = 100 * correct / total
213     print('Epoch [{}/{}], Validation Accuracy: {:.2f}%
214         '.format(epoch+1, num_epochs, accuracy))
215
216 # Problem 4
217
218 # Create data loaders for batch training
219 train_dataset = torchvision.datasets.MNIST(root='./
220 data', train=True, download=True, transform=
221     transform)
222 test_dataset = torchvision.datasets.MNIST(root='./data
223 ', train=False, download=True, transform=transform)
224 valid_dataset = torchvision.datasets.MNIST(root='./
225 data', train=True, download=True, transform=
226     transform)
227
228 train_loader = torch.utils.data.DataLoader(
229     train_dataset, batch_size=64, shuffle=True)

```

```

221 valid_loader = torch.utils.data.DataLoader(
    valid_dataset, batch_size=64, shuffle=True)
222 test_loader = torch.utils.data.DataLoader(test_dataset
    , batch_size=64, shuffle=False)
223
224 # Define the network structure
225 class ConvNet(nn.Module):
226     def __init__(self):
227         super(ConvNet, self).__init__()
228         # First convolutional layer
229         self.conv1 = nn.Conv2d(in_channels=1,
            out_channels=32, kernel_size=3, stride=1,
            padding=1)
230         # Second convolutional layer
231         self.conv2 = nn.Conv2d(in_channels=32,
            out_channels=64, kernel_size=3, stride=1,
            padding=1)
232         # Fully connected layer
233         self.fc = nn.Linear(in_features=64 * 7 * 7,
            out_features=10)
234         # Pooling layer
235         self.pool = nn.MaxPool2d(kernel_size=2, stride
            =2)
236         # Activation function
237         self.relu = nn.ReLU()
238
239     def forward(self, x):
240         # Apply first convolutional layer
241         x = self.relu(self.conv1(x))
242         x = self.pool(x)
243         # Apply second convolutional layer
244         x = self.relu(self.conv2(x))
245         x = self.pool(x)
246         # Flatten the output for the fully connected
            layer
247         x = x.view(-1, 64 * 7 * 7)
248         x = self.fc(x)
249         return x
250
251 # Initialize the network
252 model = ConvNet()
253
254 # Loss function
255 criterion = nn.CrossEntropyLoss()
256
257 # Optimizer
258 optimizer = optim.SGD(model.parameters(), lr=0.01,
    momentum=0.9, weight_decay=weight_decay)
259
260 # Train the network

```

```

261 for epoch in range(40): # loop over the dataset
    multiple times
262     running_loss = 0.0
263     for i, data in enumerate(train_loader, 0):
264         # get the inputs; data is a list of [inputs,
            labels]
265         inputs, labels = data
266
267         # zero the parameter gradients
268         optimizer.zero_grad()
269
270         # forward + backward + optimize
271         outputs = model(inputs)
272         loss = criterion(outputs, labels)
273         loss.backward()
274         optimizer.step()
275
276         running_loss += loss.item()
277
278     # Validation accuracy
279     correct = 0
280     total = 0
281     with torch.no_grad():
282         for data in valid_loader:
283             images, labels = data
284             outputs = model(images)
285             _, predicted = torch.max(outputs.data, 1)
286             total += labels.size(0)
287             correct += (predicted == labels).sum().
                item()
288
289     # Print statistics
290     print(f'Epoch_{epoch+1}, Loss:_{running_loss/_
        len(train_loader)},'
291           f'Validation_Accuracy:_{100*_correct/_
        total}%')
292
293     print('Finished_Training')

```